

Chapter 4 Classification

Austin R.J. Downey


Classification

Machine learning tasks often fall into two categories:

- ▶ Regression
 - ▶ Predict continuous values
- ▶ Classification
 - ▶ Predict discrete classes

In this chapter we focus on:

- ▶ Binary classification
- ▶ Multiclass classification
- ▶ Performance metrics



Linear decision boundaries separating two classes.

Binary Classifier


A binary classifier assigns inputs to one of two classes.

Decision rule:

$$\hat{y} = \begin{cases} 1 & \text{if } \boldsymbol{\theta}^\top \mathbf{x} + b \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

- ▶ $\mathbf{x}$: input features
- ▶ $\boldsymbol{\theta}$: model weights
- ▶ b : bias term

Training adjusts $\boldsymbol{\theta}$ and b to minimize classification error.



Multiple linear decision boundaries can correctly separate the data.

Data: MNIST Dataset

The MNIST dataset contains:

- ▶ 70,000 handwritten digits
- ▶ Classes 0 through 9
- ▶ Standard benchmark for classification

Typical split:

- ▶ 60,000 training images
- ▶ 10,000 testing images



Examples of MNIST digits.

Data: What is a Feature?

Each MNIST image is:

$$28 \times 28 = 784$$


pixels.

In machine learning:

- ▶ Each pixel is a feature
- ▶ Pixel intensity ranges from 0 to 255
- ▶ One image becomes a vector with 784 features

For MNIST:

$$\mathbf{x} = [x_1, x_2, \dots, x_{784}]$$



Each pixel becomes a feature used by the classifier.

Example 4.1: MNIST Dataset

To simplify the classification problem, we begin by building a: **5-detector**

Goal:

- ▶ Determine whether a digit is a 5
- ▶ Binary output:
 - ▶ 1 = digit is a 5
 - ▶ 0 = digit is not a 5

Why choose 5s?

- ▶ One of the most difficult digits to classify
- ▶ Frequently confused with 3s and 8s
- ▶ Good benchmark for binary classification

This example demonstrates how to use:

```
sklearn.datasets.fetch_openml
```

to load and visualize the MNIST dataset.

Regularized Linear Classifier

A practical way to build the MNIST “5-detector” is with a regularized linear classifier trained using Stochastic Gradient Descent.

- ▶ Efficient on large datasets
- ▶ Simple to implement
- ▶ Updates one example at a time

Main drawbacks:

- ▶ Needs careful hyperparameter tuning
- ▶ Sensitive to feature scaling

Training Objective

Given labeled samples $(\mathbf{x}^{(i)}, y^{(i)})$ with $y^{(i)} \in \{0, 1\}$, the “5-detector” learns parameters by minimizing:

$$J(\boldsymbol{\theta}) = \frac{1}{m} \sum_{i=1}^m \max(0, 1 - y^{(i)} \boldsymbol{\theta}^T \mathbf{x}^{(i)}) + \frac{\lambda}{2} \|\boldsymbol{\theta}_1\|_2^2$$

- ▶ First term: hinge loss
- ▶ Second term: ℓ_2 regularization
- ▶ Bias term is not regularized

Note

In scikit-learn, this model is implemented as `SGDClassifier`.

Example 4.2: Regularized Linear Classifier

- ▶ Train a regularized linear classifier with `SGDClassifier`
- ▶ Use the MNIST dataset to distinguish the digit “5” from all others
- ▶ Show how regularization improves generalization
- ▶ Demonstrate prediction on a selected handwritten digit

This example illustrates how a sparse linear decision boundary can work well for binary digit classification.

Performance Measures for Classification

Evaluating a classifier is often more challenging than evaluating a regressor.

Questions we might ask:

- ▶ How many predictions were correct?
- ▶ What kinds of mistakes were made?
- ▶ Which classes are most often confused?

A common tool for answering these questions is the:

Confusion Matrix

Binary Confusion Matrix

Two common error types:

- ▶ Type I Error:
 - ▶ False Positive (FP)
- ▶ Type II Error:
 - ▶ False Negative (FN)

A confusion matrix summarizes:

- ▶ True Positives (TP)
- ▶ True Negatives (TN)
- ▶ False Positives (FP)
- ▶ False Negatives (FN)

		Predicted	
		negative	positive
Actual	negative	3 7 9	8 3
	positive	5 5	5 5

Binary confusion matrix.

Confusion Matrix

Rows correspond to:

- Actual classes


Columns correspond to:

- Predicted classes

A good classifier has:

- Large values on the diagonal
- Small values elsewhere

Off-diagonal entries indicate misclassifications.



MNIST confusion matrix.

Confusion Matrix Terminology

A confusion matrix summarizes classifier predictions.

- ▶ Rows = actual class
- ▶ Columns = predicted class
- ▶ TP: True Positive
- ▶ TN: True Negative
- ▶ FP: False Positive
- ▶ FN: False Negative

A perfect classifier contains values only along the main diagonal.

TN	FP	TN
FN	TP	FN
TN	FP	TN

Example 4.3: Binary Confusion Matrix

- ▶ Train a binary classifier to detect the digit “5”
- ▶ Generate a confusion matrix
- ▶ Compute:
 - ▶ True Positives (TP)
 - ▶ True Negatives (TN)
 - ▶ False Positives (FP)
 - ▶ False Negatives (FN)
- ▶ Examine examples of:
 - ▶ Missed 5s (false negatives)
 - ▶ Incorrectly detected 5s (false positives)
- ▶ Use the confusion matrix to understand where the classifier fails

Accuracy

Accuracy is defined as

$$\text{Accuracy} = \frac{\text{Correct Predictions}}{\text{Total Predictions}}$$

or

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

- ▶ Simple and intuitive
- ▶ Often the first metric reported
- ▶ Does not distinguish between error types

Why Accuracy Isn't Enough

A classifier may achieve high accuracy while still making costly mistakes.

Examples:

- ▶ Credit card fraud detection
- ▶ Medical diagnosis
- ▶ Network intrusion detection

A false negative may be much more costly than a false positive.

We therefore introduce:

Precision Recall


Selected and Relevant Instances

Definitions:

$$\text{Selected Instances} = TP + FP$$

$$\text{Relevant Instances} = TP + FN$$

These quantities form the basis of precision and recall.



how many selected instances are relevant?

$$\text{Precision} = \frac{\text{orange semi-circle}}{\text{orange and blue semi-circle}}$$

how many relevant instances are selected?

$$\text{Recall} = \frac{\text{orange semi-circle}}{\text{orange rectangle}}$$

Precision

Precision answers:

Of all positive predictions, how many were correct?

$$\text{Precision} = \frac{TP}{TP + FP}$$

High precision implies:

- ▶ Few false positives
- ▶ Positive predictions are trustworthy

Recall

Recall answers:

Of all actual positives, how many were found?

$$\text{Recall} = \frac{TP}{TP + FN}$$

High recall implies:

- ▶ Few false negatives
- ▶ Important cases are not missed

F1 Score

The F1 score balances precision and recall.

$$F_1 = 2 \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

Equivalent form:

$$F_1 = \frac{TP}{TP + \frac{FP+FN}{2}}$$

Useful when both precision and recall are important.


Precision–Recall Tradeoff

Changing the classification threshold changes:

- Precision
- Recall
- F1 score

Higher threshold:

- Higher precision
- Lower recall




Choosing a Threshold

Threshold selection depends on the application.

Examples:

- ▶ Medical diagnosis → high recall
- ▶ Spam detection → high precision

The F1 score often helps identify a good compromise.




Precision–Recall Curve

The precision–recall curve summarizes classifier performance across all thresholds.

Better classifiers typically have:

- ▶ Higher precision
- ▶ Higher recall
- ▶ Larger AUC-PR

Applications determine the desired operating point.



Example 4.4: Accuracy, Precision, and Recall

- ▶ Train a binary classifier to detect the digit “5”
- ▶ Compute:
 - ▶ Accuracy
 - ▶ Precision
 - ▶ Recall
 - ▶ F1 Score
- ▶ Compare manual and sklearn calculations
- ▶ Visualize threshold effects
- ▶ Explore precision–recall tradeoffs

k -fold Cross-Validation

Training an SGD classifier does not always produce identical results.

Reasons:

- ▶ Random initialization
- ▶ Random ordering of samples
- ▶ Stochastic optimization

Possible solutions:

- ▶ More training epochs
- ▶ More data
- ▶ Cross-validation

Goal:

- ▶ Obtain a more reliable estimate of model performance


Performance Variability

Repeated training can produce different results.

Observations:

- Accuracy varies
- Precision varies
- Recall varies

Cross-validation reduces this variability and produces more stable estimates.




How k -fold Cross-Validation Works

Procedure:

1. Split data into k folds
2. Train on $k - 1$ folds
3. Validate on the remaining fold
4. Repeat for all folds
5. Average the results

Advantages:

- Uses all available data
- Better estimate of performance
- Reduces dependence on a single train/test split



Example 4.5: k -fold Cross-Validation

- ▶ Train a binary classifier to detect the digit “5”
- ▶ Repeatedly train an SGD classifier
- ▶ Observe variability in:
 - ▶ Accuracy
 - ▶ Precision
 - ▶ Recall
- ▶ Apply:

```
sk.model_selection.cross_val_predict
```

- ▶ Demonstrate how k -fold cross-validation provides more reliable performance estimates

Multiclass Classification

Binary classifiers distinguish between:

2 classes

Multiclass classifiers distinguish between:

$N > 2$ classes

For MNIST:

10 classes

corresponding to the digits:

$0, 1, 2, \dots, 9$

A common strategy is to combine multiple binary classifiers.

One-vs-Rest (OvR)

One-vs-Rest trains:


N

binary classifiers.

For MNIST:

- ▶ 0-detector
- ▶ 1-detector
- ▶ 2-detector
- ▶ ...
- ▶ 9-detector

The class with the highest score is selected.



One-vs-One (OvO)

One-vs-One trains a classifier for every pair of classes.

Number of classifiers:


$$\frac{N(N-1)}{2}$$

For MNIST:

$$\frac{10(10-1)}{2} = 45$$

classifiers.

The class with the most wins is selected.



Example 4.6: Multiclass Regularized Linear Classifier

- ▶ Extend `SGDClassifier` to multiclass classification
- ▶ Compare:
 - ▶ One-vs-Rest (OvR)
 - ▶ One-vs-One (OvO)
 - ▶ Scikit-Learn automatic strategy
- ▶ Train classifiers on the MNIST dataset
- ▶ Measure computational cost and performance
- ▶ Compare strategies for distinguishing the 10 digit classes

Performance Measures for Multiclass Classification

Many of the metrics used for binary classification extend naturally to multiclass problems.

Common evaluation tools include:

- ▶ Confusion matrices
- ▶ Accuracy
- ▶ Precision and recall
- ▶ Error analysis

For multiclass problems, the confusion matrix is often the most informative starting point.

Multiclass Confusion Matrix

Rows correspond to:

- Actual digit


Columns correspond to:

- Predicted digit

A good classifier has:

- Large values on the diagonal
- Small values elsewhere

The MNIST classifier performs well because most entries lie on the diagonal.




Normalized Error Matrix

To better visualize mistakes:

- ▶ Normalize each row
- ▶ Remove the diagonal entries
- ▶ Display only classification errors

This highlights which digits are most often confused.

Brighter regions correspond to larger error rates.



Analyzing Individual Errors

The classifier frequently confuses:

$$3 \leftrightarrow 5$$

Reasons:

- ▶ Similar handwritten shapes
- ▶ Small stroke differences
- ▶ Sensitivity to translation and rotation

Improvement strategies:

- ▶ Better preprocessing
- ▶ Additional features
- ▶ Nonlinear models

digits classified as a 3					digits classified as a 5				
correctly classified as 3s	3	3	3	3	3	3	3	3	incorrectly classified as 5s
	3	3	3	3	3	3	3	3	
	3	3	3	3	3	3	3	3	
	3	3	3	3	3	3	3	3	
	3	3	3	3	3	3	3	3	
incorrectly classified as 3s	5	5	5	5	5	5	5	5	correctly classified as 5s
	5	5	5	5	5	5	5	5	
	5	5	5	5	5	5	5	5	
	5	5	5	5	5	5	5	5	
	5	5	5	5	5	5	5	5	

Example 4.7: Multiclass Confusion Matrix

- ▶ Train a One-vs-One classifier using SGD
- ▶ Perform 3-fold cross-validation
- ▶ Construct a multiclass confusion matrix
- ▶ Visualize:
 - ▶ Raw classification counts
 - ▶ Normalized classification errors
- ▶ Identify commonly confused digits
- ▶ Analyze classifier weaknesses
- ▶ Use the results to guide model improvements