

```

1  """
2  Example 2.1 Linear Regression
3  @author: Austin Downey
4  """
5
6  import IPython as IP
7  IP.get_ipython().run_line_magic('reset', '-sf')
8
9  import numpy as np
10 import matplotlib.pyplot as plt
11 import sklearn as sk
12 from sklearn import datasets
13 from sklearn.linear_model import LinearRegression
14
15 plt.close('all')
16
17 %% load data
18
19 ames = sk.datasets.fetch_openml(name="house_prices", as_frame=True, parser='auto')
20 target = ames['target'].values
21 data = ames['data']
22 YrSold = data['YrSold'].values # Year Sold (YYYY)
23 MoSold = data['MoSold'].values # Month Sold (MM)
24 OverallCond = data['OverallCond'].values # OverallCond: Rates the overall condition of
the house
25 GrLivArea = data['GrLivArea'].values # Above grade (ground) living area square feet
26 BedroomAbvGr = data['BedroomAbvGr'].values # Bedrooms above grade (does NOT include
basement bedrooms)
27
28 # Ask a home buyer to describe their dream house, and they probably won't begin
29 # with the height of the basement ceiling or the proximity to an east-west railroad.
30 # But this playground competition's dataset proves that much more influences price
31 # negotiations than the number of bedrooms or a white-picket fence.
32
33 # With 79 explanatory variables describing (almost) every aspect of residential
34 # homes in Ames, Iowa, this competition challenges you to predict the final price
35 # of each home.
36
37 # Plot a few of the interesting features vs the target (price). In particular,
38 # let's plot the number of rooms vs. the price.
39 plt.figure()
40 plt.plot(GrLivArea, target, 'o', markersize=2)
41 plt.xlabel('Above grade (ground) living area square feet')
42 plt.ylabel('price (USD)')
43 plt.grid(True)
44 plt.tight_layout()
45
46 %% Build a model for the data
47 X = GrLivArea
48 Y = target
49 model_X = np.linspace(0, 5000)
50
51 theta_1 = 0
52 theta_2 = 100
53 model_Y_manual = theta_1 + theta_2 * model_X
54
55 plt.figure()
56 plt.plot(X, Y, 'o', markersize=2, label='data')
57 plt.plot(model_X, model_Y_manual, '--', label='manual fit')
58 plt.xlabel('Above grade (ground) living area square feet')
59 plt.ylabel('price (USD)')
60 plt.grid(True)
61 #plt.xlim([3.5, 9])
62 #plt.ylim([0, 50000])
63 plt.legend(framealpha=1)
64 plt.tight_layout()
65

```

```

66 # add a dimension to the data as math is easier in 2d arrays and sk learn only
67 # takes 2d arrays
68 X = np.expand_dims(X,axis=1)
69 Y = np.expand_dims(Y,axis=1)
70 model_X = np.expand_dims(model_X,axis=1)
71
72 %% compute the linear regression solution using the closed form solution
73
74 # compute
75 X_b = np.ones((X.shape[0],2))
76 X_b[:,1] = X.T # add x0 = 1 to each instance
77
78 theta_closed_form = np.linalg.inv(X_b.T@X_b)@X_b.T@Y
79
80 model_y_closed_form = theta_closed_form[0] + theta_closed_form[1]*3000
81 model_Y_closed_form = theta_closed_form[0] + theta_closed_form[1]*model_X
82
83 plt.figure()
84 plt.plot(X,Y,'o',markersize=3,label='data points')
85 plt.xlabel('Above grade (ground) living area square feet')
86 plt.ylabel('price (USD)')
87 plt.plot(3000,model_y_closed_form,'dr',markersize=10,zorder=10,
88         label='inferred data point')
89 plt.plot(model_X,model_Y_closed_form,'-',label='normal equation')
90 plt.grid(True)
91 plt.legend()
92 plt.tight_layout()
93
94 %% compute the linear regression solution using gradient descent
95
96 eta = 0.00000001 # learning rate
97 n_iterations = 100
98 m = X.shape[0]
99 theta_gradient_descent = np.random.randn(2,1) # random initialization
100 for iteration in range(n_iterations):
101     gradients = 2/m * X_b.T.dot(X_b.dot(theta_gradient_descent) - Y)
102     theta_gradient_descent = theta_gradient_descent - eta * gradients
103
104 print(theta_gradient_descent)
105
106 model_Y_gradient_descent = theta_gradient_descent[0] \
107     + theta_gradient_descent[1]*model_X
108
109 plt.figure()
110 plt.plot(X,Y,'o',markersize=3,label='data points')
111 plt.xlabel('Above grade (ground) living area square feet')
112 plt.ylabel('price (USD)')
113 plt.plot(model_X,model_Y_closed_form,'-',label='normal equation')
114 plt.plot(model_X,model_Y_gradient_descent,':',label='gradient descent')
115 plt.grid(True)
116 plt.legend()
117 plt.tight_layout()
118
119 %% compute the linear regression solution using sk-learn
120
121 # build and train a closed from linear regression model in sk-learn
122 model_LR = sk.linear_model.LinearRegression()
123 model_LR.fit(X,Y[:,0])
124 model_Y_sk_LR = model_LR.predict(model_X)
125
126 # build and train a Stochastic Gradient Descent linear regression model in sk-learn.
127 # Note that in running the model, the best way to do this would be to use a pipeline
128 # =with feature scaling. However, here we just set 'eta0' to a low value, this
129 # is done only for educational # purposes and is not the ideal methodology in
130 # terms of system robustness.
131 model_SGD = sk.linear_model.SGDRegressor(learning_rate='constant',eta0=0.00000001)
132 model_SGD.fit(X,Y[:,0])

```

```
133 model_Y_sk_SGD = model_SGD.predict(model_X)
134
135 # plot the modeled results
136 plt.figure()
137 plt.plot(X,Y,'o',markersize=2,label='data')
138 plt.plot(model_X,model_Y_closed_form,'-',label='normal equation')
139 plt.plot(model_X,model_Y_gradient_descent,'--',label='gradient descent')
140 plt.plot(model_X,model_Y_sk_LR,':',label='sklearn normal equation')
141 plt.plot(model_X,model_Y_sk_SGD,'-.',label='sklearn stochastic gradient descent')
142
143 plt.xlabel('Above grade (ground) living area square feet')
144 plt.ylabel('price (USD)')
145 plt.grid(True)
146 plt.legend(framealpha=1)
147 plt.tight_layout()
```